Upskilling Execution for Microsoft Fabric

Upskilling Execution for Microsoft Fabric

Workshop

Duration: 1-3 Days   [Remote / Onsite]
Difficulty Level: 300 - Advanced

Description

This engagement is designed to help build and deliver your capability and skilling plan that includes the knowledge you need to achieve successful and insightful analytics with Microsoft Fabric. Microsoft Fabric is a comprehensive analytics platform that unifies all the necessary data and analytics tools for organizations. Microsoft Fabric empowers both data and business professionals to unlock the potential of data and pave the way for the AI with the integration of technologies such as Azure Data Factory, Azure Synapse Analytics, and Power BI.

Objectives

- Learn the basics with an overall introduction experience applicable to all roles.
- Enable choice in providing role and technology focus areas for deeper discussions.

Outcomes

- Introduction – Overall view of Microsoft Fabric capabilities.
- Deep Dives – Specific insights to modules of Microsoft Fabric.

Methodology

Learn by example: Learn with presentations, demonstrations and group discussion.

Requirements

Participants
Participation of key roles is necessary for the overall success and delivery of this engagement.

- AI Engineer
- Data Analyst
- Database Administrator
- Data Engineer
- Data Scientist
- Cloud Administrator
- Solution Architect

Skill Requirements

- Basic knowledge of the Microsoft Azure platform and Azure portal.

Requirements

- Lab experience with Microsoft Trial tenant will be provided. Microsoft Trial tenant restrictions may apply.

Microsoft Fabric (Day 1: 4/21)

Topics	Time (Eastern Time)	Contents	Owner
Product Overview • SaaS Architecture • OneLake in Details • Security & Governance Overview	10:00-11:00 (60)	Overview of Microsoft Fabric's unified SaaS architecture, including OneLake for centralized storage and integrated governance and security capabilities. Learn how Fabric supports Power BI, Synapse, and Data Factory workloads.	Rodrigo Gomez Jim Ziegler
Workloads Overview – Data Integration	11:00-11:10 (10)	Introduction to Microsoft Fabric's Data Integration workload. Understand the core capabilities, connectors, and how it supports building data pipelines and ingestion workflows.	Shingo Sakamoto
Lab 1: Data Prep	11:10-11:50 (40)	Hands-on lab to build a data pipeline and transformation dataflow using Dataflows Gen2. Ingest data from SharePoint into Microsoft Fabric's lakehouse using Power Query and transformations.	Shingo Sakamoto Yuka Kato
Lunch Break	11:50-12:50 (60)		-
Workloads Overview – Lakehouse	12:50-13:00 (10)	Learn the lakehouse architecture in Microsoft Fabric. Understand how Delta Lake, notebooks, and tables enable scalable analytics with built-in security and workspace-level control.	Shingo Sakamoto
Lab 2: Data Lakehouse & Warehouse & Report	13:00-14:20 (80)	Use a Spark notebook to analyze files stored in the lakehouse. Learn how to create a Delta Table, connect to the warehouse, and build a Power BI report to visualize insights.	Shingo Sakamoto Yuka Kato
Q&A	14:20-14:30 (10)	Clarify concepts around lakehouse architecture, OneLake, SaaS integration, and governance. Understand how Microsoft Fabric delivers analytics flexibility through a unified platform. Discuss the next steps.	ALL

Data Science in Microsoft Fabric (Day 2: 4/23)

Topics	Time (Eastern Time)	Contents	Owner
Use Case	10:00-10:40 (40)	Introduce a real-world business challenge to solve using data analytics serices in Microsoft Fabric. Walk through goals, KPIs, and expected outcomes. • Audit • PIDSA • PCEC • PEXNA IT	Ken/Tina (Audit) Jerry Deng (PIDSA) Brian Rowe (PCEC) Brendan Byrne (PEXNA)
Lesson 1: Data Science in MS Fabric	10:40-10:50 (10)	Explore the integrated data science capabilities within Microsoft Fabric, including how data engineers and scientists can collaborate using a unified platform. Learn how Fabric supports the end-to-end machine learning lifecycle with scalable compute, built-in notebooks, and seamless data access.	Shingo Sakamoto
Lesson 2: Model Training & Tracking	10:50-11:00 (10)	Understand how to train and track models within Microsoft Fabric using MLflow. Learn how experiments, parameters, and metrics are automatically logged, making it easier to monitor performance, reproduce results, and manage model versions.	Shingo Sakamoto
Lab 1: Get started with data science	11:00-12:00 (60)	Ingest data, explore the data in a notebook, process the data with the Data Wrangler, and train two types of models. By performing all these steps, you'll be able to explore the data science features in Microsoft Fabric.	Shingo Sakamoto Yuka Kato
Lunch Break	12:00-13:00 (60)		
Lesson 3: Azure OpenAI in Fabric, Copilot and Data Agent	13:00-13:30 (30)	Explore how Azure OpenAI can be integrated into Microsoft Fabric to build intelligent, language-powered applications. Learn how to use large language models (LLMs) for tasks like data summarization, classification, and natural language querying—all within the Fabric environment.	Shingo Sakamoto
Q&A	13:30-14:00 (30)	Open floor for questions and discussion. Address common challenges, best practices, and use case extensions. Discuss the next steps. Share learning resources and encourage hands-on exploration.	ALL